

A Comparative Analysis of Multi-Level Generative AI Policies in NYC and Hong Kong
Higher Education Landscape

Yan Mak, Emy Yamamoto and Yingxuan Yan

Teachers College, Columbia University

HUDK 4051: Learning Analytics: Process and Theory

Dr. Renzhe Yu

May 16, 2026

Introduction

Generative artificial intelligence (GenAI) has rapidly become a central governance issue in higher education because it affects how universities define learning, academic integrity and responsible technology use. Since the public launch of ChatGPT in November 2022 the swift proliferation of large language models (e.g. Claude, Gemini, and Copilot), governments, universities and instructors worldwide have been compelled to respond to the growing presence of these technologies in education. Prior research shows that universities in the U.S. are responding to GenAI with a mix of policies, guidelines, resources and faculty support while adopting an “open but cautious approach” that recognizes both opportunities and risks (Wang et al, 2024). At the federal level, policy has leaned toward encouraging innovation but without specific binding regulation, leaving state and institutional actors to fill the gap. For example, the 2025 White House executive order on AI education states that U.S. policy is to promote AI literacy, appropriate AI integration in education, educator training, and workforce readiness, while the U.S. Department of Education’s AI guidance frames AI use as allowable under federal education programs aligned with regulatory requirements. (The White House, 2025; U.S. Department of Education, 2025). Consequently, much of the responsibility for translating broad federal guidance into concrete GenAI policies falls to institutions and individual instructors.

While these high-level patterns are well documented, less attention has been paid to how GenAI policies are translated across the layers of governance that ultimately shape student experience. Federal and state guidance does not automatically determine institutional policy, and institutional policy does not automatically translate into the classroom. Yet, it is at the classroom level that students most directly encounter GenAI policies, often through syllabus statements or instructor-level rules.

GenAI policy in education is still taking shape unevenly across governance levels without strict enforceable measures. The OECD’s digital education work captures the tension between governance systems that are trying to unlock GenAI’s potential for learning while managing risks around academic integrity, privacy and algorithmic bias (OECD Digital Education Outlook, 2023). However most institutions still rely on non-enforceable guidance rather than formal regulation. Similarly, the Digital Promise’s review of state AI guidance shows that AI-in-education policy often remains exploratory, with many systems still developing evaluation practices and fewer offering systematic evidence about impact on learning outcomes (Ruiz et al, 2025).

This project examines GenAI policy as a multi-level governance issue in the higher education landscapes of New York City and Hong Kong. These cities were selected as the primary sites for this study because both serve as preeminent global hubs for finance and higher education. Despite these similarities, their respective national contexts offer a compelling contrast in educational governance and pedagogy: the U.S. system is characterized by its decentralized structure and emphasis on inquiry-based

learning, whereas the Chinese system features a centralized framework focused on mastery-based learning and rigorous standardized assessment (Peterson et al., 2021). Comparing them allows us to ask not only what each region says about GenAI, but also whether their policies remain coherent as they move from broad national or governmental guidance down to course-level practice.

To capture this trajectory, this study analyzes documents across five tiers of governance, from national and federal documents to course syllabi, using qualitative content coding to enable quantitative comparison. Our central landscape research question is: How do Generative AI policies manifest across the five tiers of governance (National/Federal, State/Territory, Institutional, Faculty/School, and Course-wide) in NYC and Hong Kong?

New York City and Hong Kong Landscapes

This section reviews the higher education landscapes of New York City (NYC) and Hong Kong (HK). They represent two of the world's most concentrated and influential academic ecosystems, but are governed through distinctly different structural and political frameworks. New York City is home to over 100 degree-granting institutions and a student population exceeding half a million (NYCEDC, 2024). This ecosystem is characterized by a high degree of institutional plurality, varying from large public systems such as The City University of New York (CUNY) to elite private research universities such as Columbia and New York University (NYU). Policy authority is distributed across federal, state (New York State Education Department), and institutional levels, with each layer exercising considerable autonomy (NYSED, 2025).

In contrast, Hong Kong's higher education system is distinguished by its high density of world-ranked institutions within a compact geographic area. The sector is primarily governed and funded through the University Grants Committee (UGC), which oversees eight publicly-funded universities (UGC, 2025). Despite its smaller scale, three of these — the University of Hong Kong (HKU), the Hong Kong Polytechnic University (PolyU), and Lingnan University — consistently appear in global rankings such as the QS World University Rankings (QS, 2026) (See Table 1)

Table 1. *Higher Education Landscape of Hong Kong vs. New York City at a Glance*

	<i>Hong Kong</i>	<i>New York City</i>
Total Degree-granting Institutions	22	100+

Total Student Enrollment (2024-25)	192,000+	503,000+
QS Top 100 Universities (2026)	5 (HKU, CUHK, HKUST, PolyU, CityU)	2 (Columbia, NYU)
Total QS Top 1000 Universities (2026)	9	5
Primary Funding Model	Public (UGC-funded)	Private & Public/CUNY
Regulatory Body	University Grants Committee (UGC)	New York State Education Department (NYSED)

It is important to note that Hong Kong operates under the "One Country, Two Systems" principle, under which educational and technological governance is a local responsibility rather than a direct mandate of the central government in Beijing (Basic Law, 1990, Ch. VI). Nonetheless, Hong Kong is increasingly embedded in the national development strategy of the Guangdong-Hong Kong-Macao Greater Bay Area, which positions the region as a hub for innovation and talent cultivation aligned with mainland China's broader technology policy (The Government of the Hong Kong Special Administrative Region, 2025). Within this regional framework, "Innovation and Technology" and "Education" serve as two of the major policy areas (Constitutional and Mainland Affairs Bureau, 2018). This dual identity of globally networked yet nationally embedded makes Hong Kong a distinctive case for understanding how international, national, and local pressures intersect in education policy.

Institutional Selection and Pairing

For stratified comparison, we selected the following institutions:

- **HK:** The University of Hong Kong (HKU), The Hong Kong Polytechnic University (PolyU), and Lingnan University.
- **NYC:** Columbia University, New York University (NYU), and The City University of New York (CUNY).

These institutions were paired to represent three distinct levels of the global academic hierarchy. **Columbia and HKU** serve as the "global flagships" of their respective regions, consistently ranking at the top of international tables. **NYU and PolyU** represent high-impact, urban-integrated research universities with substantial international footprints and professionally-oriented programs. **CUNY and**

Lingnan complete the comparison. CUNY as a large-scale public system serving a diverse student body, and Lingnan as a publicly-funded liberal arts university known for teaching-intensive pedagogy.

This selection also captures meaningful variance in governance type: private non-profit models (Columbia, NYU), a large public system (CUNY), and publicly-funded statutory bodies (HKU, PolyU, Lingnan). Such diversity allows for analysis of how different governance structures shape policy adoption. A final consideration was transparency: each selected institution provides publicly accessible documents regarding AI governance, academic integrity, or teaching guidance, ensuring data availability for our content analysis.

From Single-Layer to Multi-Layer Policy Analysis

Existing scholarship on GenAI in higher education has focused predominantly on single-layer analyses, most commonly examining either institutional-level policy statements (Wang et al., 2024) or national policy frameworks (Ruiz et al., 2025). Wang et al.'s coding framework, developed for analyzing university-level GenAI statements, provided foundational dimensions such as Stance, Disclosure, Discretion, and Focus that we adopt and extend in this study. However, the framework was designed to examine policies at a single layer of governance (i.e., institutional) rather than to trace coherence across multiple governance levels.

A separate strand of policy research has long emphasized that policies do not exist in isolation but operate within nested layers of governance. The concept of vertical policy coherence — the degree to which policies remain aligned as authority moves from national to regional to local levels — has been used productively in environmental governance (Sandström et al., 2020) and sustainable development (Bizikova, 2023). The European Commission's EPAL platform explicitly frames vertical coherence as asking "whether leaders within learning providers effectively communicate and implement policy all the way down to the classroom" which is a question that directly reflects the central concerns of our study.

By combining Wang et al.'s coding dimensions with the multi-tier logic of vertical coherence, this study contributes a framework for analyzing GenAI policy not as a static institutional document, but as a layered system that may attenuate, transform, or contradict itself as it moves from the national level into classroom practice.

Methods and Data

This study used a comparative multi-level policy landscape analysis to examine how Generative AI governance is communicated across higher education systems in New York City and Hong Kong. To do this, we collected publicly available documents from government websites, university resource pages, and open course syllabi.

Governance Framework

We organized documents using a five-tier framework distinguishing between external and internal policy levels (see Table 2). External governance covered national and city-level documents, while internal governance covered institutional, faculty, and course-level documents. This structure allowed us to trace how broad AI policy is translated into the rules students actually encounter in the classroom.

Table 2. Five-Tier Governance Framework

Domain	Tier #	Option	Definition	Sample Size
External	1	National/Federal	Issued by a central government or national ministry.	5 per region
	2	State/Territory	Issued by a regional, state, or territory regulatory body.	5 per region
Internal	3	Institutional	Issued by the central university administration.	5 per institution
	4	Faculty/Department	Issued by a specific Department, or Faculty.	5 per institution
	5	Course-wide	Issued by an individual instructor for a specific class.	5 per institution

Data Collection

For external tiers (Tiers 1–2), we conducted a comprehensive territorial sweep of publicly available government and regulatory documents. For internal tiers (Tiers 3–5), we used quota-based purposive sampling, collecting five documents per tier per institution across three institutions in each region.

For the data collection workflow, we used Google AI/Gemini to conduct targeted searches for relevant documents. To protect privacy and avoid cross-session data retention, all searches were conducted in temporary or incognito mode with history disabled. These tools were used solely to identify potential sources and no document was retained based on AI output alone. Each identified source was manually retrieved, reviewed, and evaluated against our inclusion and exclusion criteria before being included in the dataset.

Searches followed a structured prompt template: "Please provide documents for [Region/Institution name] Generative AI Policy for [Tier name] for Higher Education with links within 2024–2026." Where initial results were insufficient or too broad, follow-up prompts were used to narrow by institution, document type, or topic. Collection continued iteratively until the quota of five documents per tier per region or institution was met. All final documents were accessed directly through their official source URLs and verified as publicly available at the time of collection.

Inclusion and Exclusion Criteria

To be included, a document had to meet all three of the following criteria:

- Authorized directive: Issued by an official body such as a Provost office, Registry, IT department, or Teaching and Learning Center
- Explicit GenAI content: Specifically references Generative AI, large language models, ChatGPT, or AI-related updates to academic integrity in a higher education context
- Verifiable source: Hosted on an official domain (.edu, .gov, .edu.hk, .gov.hk, .gov.cn, etc.) or available as a downloadable PDF from an official portal

Documents were excluded if they met any of the following conditions:

- Secondary sources: News articles or blog posts, unless they contained verbatim official policy text with no other accessible source
- Access-restricted: Password-protected or intranet-only documents not publicly available
- General tech policy: Standard plagiarism or technology use policies that do not explicitly mention AI

- Tool collections: Simple lists of AI tools without any governing language or pedagogical guidance

In total, 5 National/Federal records were collected for NYC and 5 for HK, 5 City/Territory records were collected for NYC and 5 for HK, 15 Institutional records were collected for NYC and 15 for HK, 15 Faculty/Department records were collected for NYC and 15 for HK, and finally 15 Course-wide documents from NYC schools and 15 from HK. This yielded a total sample of 110 documents, meaning 55 per region across all five tiers of governance. (See Table 3)

Table 3. Count and Percentage of Documents Collected by Region

Category	<i>Hong Kong (n, %)</i>	<i>New York City (n, %)</i>	<i>Total</i>
Executive Order	2 (4.1%)	4 (7.3%)	6
FAQ/Resource Page	5 (9.1%)	13 (23.6%)	18
Law/Regulation	0 (0.0%)	1 (1.8%)	1
Official Guidance	19 (34.5%)	7 (12.7%)	26
Other	2 (4.1%)	2 (3.6%)	4
Syllabus	15 (27.3%)	15 (27.3%)	30
University Policy	6 (10.9%)	13 (23.6%)	19
Total	55	55	110

Coding Framework

To systematically analyze the collected documents, we developed a deductive coding framework across six dimensions, drawing on Wang et al.'s (2024) university-level GenAI policy analysis framework and extending it to accommodate the multi-tier structure of this study. Each document was coded independently along all six dimensions. Table 4 defines each dimension, its categories, and its measurement scale.

Table 4. Coding framework dimensions, category definitions, and measurement scales

Dimension	Definition	Categories	Scale
Stance	Overall orientation toward GenAI use	Prohibitive (0) — Cautious (1) — Functional (2) — Encouraging (3)	Ordinal (0–3)
Disclosure	Whether AI use must be disclosed	Required / Recommended / Conditional / Not specified	Nominal
Discretion	Who holds rule-making authority	Gov-mandated / Institution-mandated / Faculty/Dept decides / Instructor decides / Hybrid / Not specified	Nominal
Focus (Primary)	Primary thematic concern of the document	Academic Integrity / AI Literacy / Pedagogy/Learning / Privacy/Data/Ethics / Professional Readiness / Equity	Nominal
Enforcement	Enforcement mechanism referenced	Explicit penalties / Referenced to integrity code / Guidance only / None stated	Nominal
Coverage	Who or what the policy applies to	Students and Faculty / Students only / Faculty only / Researchers only/ Whole community	Nominal

Note. Stance is the primary ordinal outcome variable. All other dimensions are nominal.

The Stance dimension is the primary analytical variable and was coded on a four-point ordinal scale reflecting the document's overall orientation toward GenAI use in higher education. To enable quantitative analysis, stance categories were converted to numerical scores (0–3) preserving their ordinal rank order. Governance tier was similarly converted to a numerical variable for statistical analysis (Tier_Numeric: 1 = National/Federal through 5 = Course-wide), allowing tier to serve as both an independent variable and an ordering principle in regression and correlation analyses. The other dimensions were analyzed as nominal categories.

To support qualitative coding for these documents, a comprehensive rubric was also developed for each dimension analyzed in this study. (See Appendix A).

Data Analysis Procedure

The statistical analysis was conducted in Python and used a combination of non-parametric tests for ordinal outcomes (Stance_Numeric) and chi-square or Fisher's exact tests for nominal coding dimensions.

Stance_Numeric is an ordinal variable and cannot be assumed to be normally distributed, non-parametric tests were the primary inferential tools for stance comparisons. Therefore, Mann-Whitney U test was used for all pairwise regional comparisons to test whether one region tends to present higher stance scores than the other. Kendall's Tau was used to assess the monotonic relationship between governance tier and stance across the entire sample, providing a directional and rank-sensitive measure of the vertical coherence trend.

Across institutions and policy tiers, stance scores were compared to examine whether GenAI policy positions varied by institution, tier, or whether stance differed across groups. The Kruskal-Wallis H-test was used for heterogeneity in stance across institutions within each tier.

Regional associations for nominal coding dimensions (Document Type, Disclosure, Discretion, Focus, Enforcement, Coverage) were tested using chi-square tests of independence. Effect sizes for nominal associations were measured using Cramér's V, with benchmarks of small ≈ 0.10 , moderate ≈ 0.30 , and large ≈ 0.50 .

Tier-specific Fisher's exact tests were additionally run for the three nominal variables showing any overall significant association (Document Type, Focus, Discretion) to identify which governance levels drive the overall regional differences.

Supplementary Language Analysis: Lexical Comparison and Network Analysis

To examine the language patterns embedded in regional GenAI policy discourse and provide additional interpretive context, we conducted lexical comparison of the key-quote passages extracted from institutional and governmental documents (Tiers 1–4). Course-level syllabi (Tier 5) were excluded from this particular analysis because their language tends to reflect individual instructor voice rather than institutional positioning.

All quoted passages from the NYC dataset were aggregated into a single corpus, and the same was done for Hong Kong. Each corpus was lowercased, stripped of URLs, punctuation, and non-alphabetic characters. To prevent surface-level fragmentation of related terms, we applied a light

stemming procedure that merged regular plural forms with their singular roots (e.g., "agencies" merged with "agency"; "instructors" with "instructor"). A custom stopword list extended the standard English stopword set to remove four categories of noise: (1) generic GenAI terminology (AI, ChatGPT, GenAI), (2) search-prompt artifacts (please, provide, document, level), (3) place and institution names (NYC, Hong Kong, Columbia, HKU, etc.), and (4) common modal and connective words (may, shall, will, however).

To visually examine how the coded attributes distribute across documents and to illustrate the structural relationships among governance categories. We constructed an undirected weighted graph of all 110 coded documents. This network was designed to test the mathematical stability of our assigned coding dimensions. Each document was represented as a node, and edges were drawn between pairs of documents according to the following weighting scheme: documents sharing both region and stance received an edge of weight 2.0; documents sharing stance but not region received an edge of weight 1.0; and documents sharing region and tier but not stance received an edge of weight 0.4. This scheme prioritizes stance similarity while allowing weaker structural ties between documents at the same hierarchical level.

The network was rendered using the ForceAtlas 2 layout algorithm, a force-directed approach widely used in network visualization that treats edges as springs pulling connected nodes together while applying repulsive forces between all node pairs (Jacomy et al., 2014). Node size was scaled to weighted degree centrality, so that nodes with more numerous and stronger connections appear larger. We additionally applied greedy modularity-based community detection to identify clusters and computed a modularity score Q as a measure of community separation quality (Newman & Girvan, 2004); values of Q above 0.3 are generally interpreted as indicating meaningful community structure.

Key Trends and Findings

Trend 1: Hong Kong is systematically more permissive than NYC

A Mann-Whitney U test was conducted because Stance is an ordinal variable not assumed to be normally distributed. The test confirmed a statistically significant difference between the two regions ($U = 1044.0$, $p = 0.003$). Hong Kong documents had a median Stance score of 2.0 (Functional), while NYC documents had a median score of 1.0 (Cautious). (See Figure 1 and 2) This indicates that, across all tiers and document types, Hong Kong governance documents are meaningfully more permissive and encouraging in their orientation toward GenAI use as an educational asset compared to New York City documents that treat AI predominantly as a threat to be contained.

Figure 1. Distribution of Policy Stance Numeric Scores by Region (N = 110)





Figure 2. *Distribution of policy stance by region*



Note. HK = Hong Kong; NYC = New York City. Stance scale: 0 = Prohibitive, 1 = Cautious, 2 = Functional, 3 = Encouraging.


Trend 2: Permissiveness erodes as governance moves to the classroom

A central theoretical concern of this study is vertical policy coherence, the degree to which policies remain aligned as they move from broad national guidance down through institutional, departmental, and classroom levels. Two questions structure this section: (1) Is there a systematic directional trend in stance as governance becomes more granular? And (2) Do the two regions follow the same pattern of change across tiers?

Kendall's Tau correlation between governance tier (Tier_Numeric, 1–5) and policy stance was -0.292 ($p < 0.001$). This significant negative monotonic association indicates that as policies move from the national/federal level toward course-level documents, they tend to become progressively more restrictive. This also suggests a weakening of vertical policy coherence: broad, higher-level support for GenAI is not fully preserved as policy is translated into institutional, departmental, and course-level policies.

For each region–tier combination, we calculated the average coded stance score across all documents in that cell and plotted the results across the five governance tiers. Figure 3 shows that both regions become less permissive as policy moves closer to the classroom, but the trajectories differ. Hong Kong begins at the maximum stance score of 3.00 at the National/Federal and City/Territory levels, then declines to 2.00 at the Institutional level and stabilizes at 1.87 at the Faculty/Department and Course-wide levels. Although Hong Kong's stance becomes less encouraging at lower governance tiers, it remains close to the Functional range even at the course level.

NYC follows a more fragmented trajectory. Its National/Federal mean begins at 2.40, drops sharply to 1.00 at the City/Territory level, rebounds to 2.07 at the Institutional level, and then declines through the Faculty/Department level before reaching 0.93 at the Course-wide level. This suggests that NYC course-level documents are closer to Prohibitive than Functional. The regional gap is largest at the City/Territory and Course-wide levels, and Hong Kong shows gradual attenuation while NYC shows a more uneven pattern.

Figure 3. Mean GenAI policy stance across five governance tiers in New York City and Hong Kong

Note. (0 = Prohibitive, 3 = Encouraging).

Trend 3: The two systems converge at the institutional level, then diverge again

Tier-by-tier Mann-Whitney U tests revealed that the regional gap is not uniformly distributed across governance levels and confirm this pattern statistically (Table 5): significant regional differences emerge at Tier 2 (City/Territory; $p = 0.004$) and Tier 5 (Course-wide; $p = 0.018$), while Tiers 1, 3, and 4 show no statistically significant difference between the two regions.

Table 5. Median policy stance by governance tier and region.

Tier	Level	HK Median (Label)	NYC Median (Label)	<i>p</i>
Tier 1	National/ Federal	3.0 (Encouraging)	3.0 (Encouraging)	0.317
Tier 2	City / Territory	3.0 (Encouraging)	1.0 (Cautious)	0.004 *
Tier 3	Institutional	2.0 (Functional)	2.0 (Functional)	0.859

Tier 4	Faculty / Department	2.0 (Functional)	2.0 (Functional)	0.266
Tier 5	Course-wide	2.0 (Functional)	0.0 (Prohibitive)	0.018 *
All tiers		2.0 (Functional)	1.0 (Cautious)	0.003 *

Note: * Significance based on Mann-Whitney U tests. Stance scale: 0 = Prohibitive, 1 = Cautious, 2 = Functional, 3 = Encouraging.

This pattern suggests that our HK and NYC sample diverge at the city/territory governance level, where regional regulatory culture most directly shapes institutional expectations. This divergence re-emerges at the course level, where individual instructor discretion plays the greatest role (See Figure 4).

Figure 4. Median policy stance across five governance tiers in New York City and Hong Kong





Figure 5 also suggests a similar finding, showing that the regional gap is strongest at the City/Territory and Course-wide levels. At the City/Territory level, 100% of Hong Kong documents are Encouraging, while 100% of NYC documents are Cautious. At the Institutional level, the regions converge, with both showing 40% Encouraging documents. However, they diverge again at the

Course-wide level: Hong Kong documents are mostly Functional (60%), while NYC documents are mostly Prohibitive (53%).

Figure 5. Policy stance distribution by governance tier and region.



Note. Each bar = 100% of documents in that tier/region. Purple = Encouraging, teal = Functional, green = Cautious, red = Prohibitive. Note the dominance of Encouraging in HK Tier 2 vs. Prohibitive in NYC Tier 5.

Trend 4: NYC shows greater within-city institutional variation


Beyond the regional pattern, we examined whether individual universities within each city followed or diverged from their region's overall stance trajectory. This question targets the Tier 3–5 range, where institutional autonomy most directly shapes policy outcomes.

Figure 6 shows the Policy Stance by Institution. In Hong Kong, policy distributions remain highly consistent across all three universities. In contrast, New York City displays substantial institutional variance. This is most evident at New York University (NYU), where 40% of the sampled documents are Prohibitive, a stark contrast to the more moderate distributions at Columbia University (13% Prohibitive) and CUNY (7% Prohibitive).

Kruskal Wallis H-Test comparing institutions within each tier revealed further differences between the two systems. It reveals that within Hong Kong, there are no significant institutional differences (between HKU, Poly U and Lingnan) at any tier in terms of median stance. All three institutions move in a parallel.

However, within NYC, there is a significant difference among institutions for Tier 4 (Faculty/Department) policies ($p = 0.039$), meaning some NYC institutions have a different median stance than others at this specific tier. (Tier 5 is borderline significant with $p = 0.050$.)

Figure 6. Policy Stance by Institution



Trend 5: Thematic priorities differ meaningfully by region

A Chi-square test of independence found a significant association between region and primary policy focus ($\chi^2 = 10.16$, $p = 0.023$, Cramér's $V = 0.316$), suggesting that thematic priorities in GenAI governance documents differ meaningfully between Hong Kong and New York City. Table 6 presents the full distribution of primary focus themes across regions.

Table 6. *Primary Focus Theme by Region*

Primary Focus Theme	Hong Kong <i>n</i> (%)	NYC <i>n</i> (%)	Total <i>n</i> (%)
Academic Integrity	25 (45.5%)	14 (25.5%)	39 (35.5%)
Pedagogy / Learning	13 (23.6%)	13 (23.6%)	26 (23.6%)
AI Literacy	3 (5.5%)	14 (25.5%)	17 (15.5%)
Privacy / Data / Ethics	6 (10.9%)	10 (18.2%)	16 (14.5%)
Professional Readiness	7 (12.7%)	4 (7.3%)	11 (10.0%)
Equity	1 (1.8%)	0 (0.0%)	1 (0.9%)
Total	55 (100%)	55 (100%)	110 (100%)

In Hong Kong, Academic Integrity was the dominant theme, appearing in 45.5% of HK documents compared to 25.5% of NYC documents. NYC presented a more distributed thematic landscape, with AI Literacy emerging as its most distinctive theme: 25.5% of NYC documents focused primarily on AI Literacy, compared to only 5.5% of Hong Kong documents. Privacy/Data/Ethics was also more prominent in NYC, appearing in 18.2% of NYC documents compared to 10.9% of Hong Kong documents. Pedagogy/Learning had identical prevalence in both regions, appearing in 23.6% of documents in each region, suggesting that the pedagogical implications of GenAI are recognized across both contexts.

Trend 6: NYC Delegates GenAI Discretion to Instructors More Than Hong Kong

Discretion showed a significant regional association overall ($\chi^2 = 18.42$, $p = 0.038$, Cramér's $V = 0.367$), with the difference being most pronounced at the course-wide tier (Fisher's exact test, $p = 0.022$). Table 7 represents the full distribution of the discretion categories.

Table 7. *Discretion category by region*

Discretion Category	<i>Hong Kong n (%)</i>	<i>NYC n (%)</i>	<i>Total n (%)</i>
Instructor decides	11 (20.0%)	27 (49.1%)	38 (34.5%)
Institution-mandated	13 (23.6%)	8 (14.5%)	21 (19.1%)
Government-mandated	10 (18.2%)	9 (16.4%)	19 (17.3%)
Hybrid	12 (21.8%)	8 (14.5%)	20 (18.2%)
Faculty / Department decides	3 (5.5%)	1 (1.8%)	4 (3.6%)
Not specified	6 (10.9%)	2 (3.6%)	8 (7.3%)
Total	55 (100%)	55 (100%)	110 (100%)

NYC documents were substantially more likely to delegate authority to individual instructors (49.1% of NYC documents), compared to just 20.0% of HK documents. Hong Kong more frequently applied institution-mandated (23.6%) or government-mandated (18.2%) discretion frameworks, consistent with its centralized governance structure. This finding helps explain why NYC course-level policies were found to be inconsistent: when individual instructors are left to decide their own rules, outcomes depend on their personal judgement rather than institutional mandated policies thus leading to more variation.

Trend 7: Lexical Patterns Support Distinct Regional Discourses

Figure 7 . Word cloud comparison of high-frequency terms in institutional and governmental policy documents (Tiers 1–4) from New York City (left) and Hong Kong (right).



Figure 7 presents the word cloud comparison. The two regions draw on visibly distinct vocabularies. New York City documents foreground a regulatory and instructor-centered register, with prominent terms including "instructor," "appropriate," "technology," "assignment," and "regulation." Words signaling oversight ("prohibited," "risk," "framework," "governance") appear with notable frequency. Hong Kong documents, in contrast, foreground a research- and learning-oriented register, with "research," "learning," "content," "assessment," "academic," and "teaching" dominating the cloud. Words signaling pedagogical concern ("integrity," "ethical," "responsible") appear consistently, suggesting an emphasis on capability building rather than restriction.

This lexical divergence suggests that the two regions are not simply taking different positions on the same set of issues, but are framing GenAI through distinct governance priorities: NYC's discourse centers on risk management and instructor-level regulation, whereas Hong Kong's discourse prioritizes capability building and institutional academic integrity. We note that our word clouds directly validate our earlier statistical findings. The prominence of the word "instructor" in the NYC cloud reflects the region's heavy reliance on course-level discretion (Trend 6), while words like "prohibited" visually capture its more restrictive policy stance (Trend 1). On the other hand, the prominence of the word "integrity" in the Hong Kong cloud perfectly aligns with our finding that Academic Integrity serves as the region's dominant policy focus (Trend 5).

Trend 8: Network Analysis Structurally Validates the Stance Coding Framework

Figure 8. *GenAI policy document network ($n = 110$, edges = 2117). Node size reflects weighted degree centrality; node color indicates stance community. ForceAtlas 2 layout; modularity $Q = 0.610$.*




Figure 8 presents the GenAI policy document network. The ForceAtlas 2 layout resolves the 110 documents into four distinct, spatially separated communities, returning a high modularity score of $Q = 0.610$. Because the edge-weighting algorithm was programmed to prioritize stance and region, this high modularity supports the structural integrity of the coding scheme. Mapping the data in this way provides a visual support of the regional asymmetries identified in our earlier statistical tests. The regional composition of the algorithmically sorted communities is highly uneven. The Encouraging community is heavily dominated by Hong Kong documents (65%), while the Prohibitive community exists as an isolated cluster containing nine NYC documents and zero Hong Kong documents. This pattern visually reinforces the finding that strict restriction of GenAI use is an exclusively NYC phenomenon within this sample.

Case Studies: Columbia University and the University of Hong Kong

The University of Hong Kong

The University of Hong Kong (HKU) illustrates the quantitative trend observed across Hong Kong's governance tiers. Our statistical data shows that Hong Kong's overall stance score drops from a strongly encouraging 3.00 at the external governmental tiers to a functional 1.87 at the course-wide level. This top-down trajectory is supported by the raw documents. At the National/Federal level (Tier 1), policies like the Ministry of Education's *"AI+ Education"* Action Plan actively encourage integrating AI into all elements of teaching to drive pedagogical innovation (Ministry of Education, 2026). Similarly, at the City/Territory level (Tier 2), the Digital Policy Office's Ethical AI Framework proactively guides organizations in maximizing AI's benefits while managing ethical risks (Digital Policy Office, 2025). HKU mirrors this functional baseline at the Institutional level by establishing comprehensive guidelines that range from equal access and AI literacy (Teaching and Learning Innovation Centre [TALIC], 2023) to transparency and responsible use in the classroom (TALIC, 2026).

However, as authority moves down to the most granular course level, instructors introduce highly cautious, context-specific restrictions. For instance, an HKU Geography syllabus (GEOG3422) sets a strict 30% threshold for AI detection software, warning that surpassing this could lead to failing grades or disciplinary action (Loo, 2025). A Psychology syllabus (PSYC7319) explicitly limits AI use to gathering information and checking grammar, strictly prohibiting its use for writing drafts (The University of Hong Kong, 2025).

We must read these restrictive cases alongside a "silent majority" of courses. Some course syllabi were excluded from our final dataset because of their absence of specific GenAI directives. Most of these courses implicitly default back to the more permissive Tier 3/4 baseline. Furthermore, because our methodology was restricted to publicly accessible documents, it is possible that some instructors provide specific, localized instructions through closed learning management systems that are not open to the public.

Independent evaluative accounts from the University Grants Committee (QAC, 2024) confirm a temporal shift in HKU's policy. In February 2023, HKU initially placed a brief halt on the student use of large language models for credit-bearing activities in response to emerging student concerns. However, this restrictive phase was strictly transitional. The university swiftly promoted a campus-wide debate on the implications of GenAI for teaching and learning and conducted workshops on policy development. A Task Force on GenAI was formed to draft a comprehensive policy that was officially endorsed by the Senate in June 2023. Following this endorsement, the university actively encouraged staff to respond to

GenAI in their assessment practices and established formal support structures, making advice available through the GenAI User Advisory Committee and the Teaching and Learning Innovation Centre (TALIC). This concrete timeline shows that university governance of GenAI is dynamic and rapidly evolving. Cross-sectional analyses that capture policy stances at a single, static moment provide a temporary snapshot, and an institution's regulatory framework can look different from one year to the next.

Columbia University

Columbia University illustrates the broader NYC pattern identified in this study: GenAI governance is relatively developed, but decision-making authority is distributed across institutional offices, schools, departments, and individual instructors. In the coded dataset, Columbia's internal documents showed a cautious-to-functional tendency. Its mean stance score was 1.8 at both the institutional and faculty/departments levels, suggesting that higher-level guidance generally acknowledged GenAI's usefulness while maintaining clear guardrails. However, at the course-wide level, the mean stance dropped to 1.2, indicating that classroom-facing policies were more restrictive than broader institutional guidance.

This pattern also reflects the decentralized U.S. policy landscape in which Columbia operates. At the Federal level, AI policy has generally encouraged innovation, AI literacy and responsible adoption without establishing detailed binding rules for higher education. For example, the White House executive order on AI education emphasizes AI literacy, educator training, and workforce readiness rather than prescribing specific university-level classroom rules (The White House, 2025). As a result, institutions such as Columbia must translate broad national priorities into their own policies which are then further interpreted by schools, departments, and instructors.

At the institutional level, Columbia does not present GenAI as something to be banned outright. Instead, university-wide guidance treats it as a powerful but risk-bearing technology that can support teaching, learning, research, and administrative work when used responsibly. Columbia's Office of the Provost recognizes GenAI as a tool transforming higher education while emphasizing privacy, security, and responsible exploration (Columbia University Office of the Provost, n.d.). Similarly, CUIT guidance reemphasizes that AI should assist and not replace human judgment. It also directs the community toward approved or enterprise-supported tools, especially when sensitive or institutional data is involved (Columbia University Information Technology, 2026). This suggests that Columbia's approach is not anti-AI, but rather focused on managing GenAI use within safer institutional boundaries.

At the faculty and department level, Columbia's policies become more varied. Teachers College guidance gives instructors flexibility to experiment with GenAI while asking them to set clear course-specific expectations (Teachers College, Columbia University, 2025). Columbia Law School takes a more restrictive default position, limiting GenAI use for exams, final papers, and submitted work unless an instructor provides written permission (Columbia Law School, n.d.). Columbia Business School frames GenAI more as a tool for productivity, innovation, and professional readiness (Columbia Business School, n.d.). These examples show that Columbia's GenAI governance is shaped by the needs and risks of different academic and professional contexts.

This variation becomes even more visible in course syllabi. For instance, a Columbia Psychology syllabus (PSYC UN 3620) allows students to use GenAI when it supports course learning goals, but requires them to revise, personalize, and properly cite any AI-assisted material (Columbia University Department of Psychology, 2024a). Another Psychology syllabus (PSYC UN 2235) permits limited AI use for editing, grammar, and presentation support, while emphasizing that students' ideas and written work must remain their own (Columbia University Department of Psychology, 2024b). In contrast, a Computer Science syllabus (COMS 3261) prohibits the use of GenAI tools on homework, and a Neuroscience syllabus (Psychobiology of Stress) states that AI-generated submissions are not acceptable (Columbia University Department of Computer Science, 2025; Columbia University Department of Psychology, 2024c).

Overall, Columbia shows how decentralized governance can create both flexibility and inconsistency. Higher-level guidance recognizes GenAI's educational and professional value, but schools and instructors translate that guidance in different ways. Compared with HKU's more coordinated institutional trajectory, Columbia's model is more distributed across internal governance layers. This makes Columbia a good example of the study's central argument: GenAI policy is not simply passed down from the top, but reshaped as it moves through the institutional layers that students actually experience.

Discussion

This study examined and created a synthesized comparison of how Generative AI policies manifest across a five-tier governance framework in New York City and Hong Kong. Our 110 coded documents (55 from each region, evenly distributed across the five governance tiers) provide the basis for multiple comparisons.

Our analysis revealed four major findings. First, a significant regional divergence exists: Hong Kong's governance documents demonstrate a fundamentally more functional and encouraging stance

toward GenAI compared to NYC's cautious, restriction-heavy approach. Second, we identified a vertical policy pattern: across both regions, policies became progressively more restrictive as authority moved from broad national/federal guidance down to granular course-level syllabi. Third, the regions exhibit different structures of decision-making authority and thematic focus. NYC delegates rule-making primarily to individual instructors with a focus on AI literacy, whereas Hong Kong relies on centralized, institution-mandated guidelines heavily focused on academic integrity. Finally, lexical and network analyses confirmed these structural differences, illustrating that NYC frames GenAI through a discursive lens of compliance and risk management, while Hong Kong frames it through a lens of capability building and pedagogy.

Implications and Next Steps

These findings suggest that both NYC and Hong Kong institutions should make GenAI policies more publicly accessible, especially at the course and department levels. Greater transparency would help students understand expectations and allow researchers to better evaluate how policy is implemented across governance tiers. For NYC in particular, universities need clearer baseline GenAI policies so students are not subject to highly inconsistent rules across courses. Institutions should also provide more incentives and support for instructors to engage with AI pedagogically, rather than defaulting to restriction. This could include faculty development, sample syllabus language, teaching grants, or recognition for responsible AI integration. At the policy level, broad encouragement of AI use is not enough; institutions need structures that connect high-level guidance to classroom practice.

Limitations

First, the sample size of three institutions per region — Columbia, NYU, and CUNY for New York City; HKU, PolyU, and Lingnan for Hong Kong — may not be entirely representative of the broader higher education landscapes. While these institutions were purposely selected and paired with justification to capture meaningful variance in governance types (e.g., global flagships, urban-integrated research universities, and public/liberal arts systems), they ultimately constitute a restricted sample. Consequently, the findings may not fully capture the policy nuances of smaller, specialised, or differently governed institutions within these two highly concentrated academic hubs.

Second, the data collection relied exclusively on publicly accessible documents. Based on our exclusion criteria, internal administrative directives, password-protected materials, and intranet-only policies were deliberately omitted. As noted in our discussion of the "silent majority," it is highly probable that many instructors provide specific, localized GenAI instructions through closed internal

systems that are inaccessible to external researchers. Thus, our dataset may underrepresent the true extent and granularity of course-level governance.

Third, the documents analyzed at Tier 4 (Faculty/Department) and Tier 5 (Course-wide) are not exhaustive. Although we implemented a structured, quota-based sampling strategy of five documents per tier per institution and meticulously documented our search protocols to ensure methodological repeatability, we cannot claim to have reviewed every existing faculty policy or course syllabus. Furthermore, the document retrieval process inherently depended on the algorithmic curation of our chosen search engine (Google AI/Gemini). Even with strict precautions including utilizing temporary or incognito browsing modes with history disabled, search engine algorithms inevitably influence which specific documents surface during targeted searches.

Finally, the qualitative content analysis introduces inherent limitations regarding coding reliability across raters. Assessing qualitative dimensions such as policy "stance," "discretion," or "primary focus" involves subjective interpretation, particularly when analyzing documents across distinctly different cultural and regulatory contexts. As noted in our analysis, the inter-rater reliability (measured via Cohen's kappa) reflects this interpretive variance, underscoring the ongoing challenge of maintaining perfect coding consistency in cross-regional, multi-layer policy analyses.

Directions for Future Research

Expanding Institutional and Regional Coverage

Future research should expand data coverage both within and beyond the current sample. Conducting an exhaustive, institution-wide analysis of Tier 4 and Tier 5 documents across the six studied universities would provide a more robust measure of the "silent majority" and allow for discipline-specific comparisons. Furthermore, broadening the sample to include more of the 22 degree-granting institutions in Hong Kong and over 100 in New York City would test whether these top-down trajectories hold true across smaller, specialized, or vocational colleges.

Longitudinal Tracking of Policy Evolution

Because university governance of GenAI is highly dynamic, cross-sectional analyses provide only a temporary snapshot. As demonstrated by HKU's rapid transition from a brief halt on GenAI in February 2023 to a comprehensive integration policy by June 2023 (QAC, 2024), regulatory frameworks evolve quickly. Longitudinally tracking this dataset from 2023 to 2026 and beyond will quantitatively measure how policies mature, revealing whether restrictive course-level syllabi eventually become more encouraging as instructors adapt to the technology.

Shifting from Policy Text to Lived Experience

Finally, future studies must close the gap between policy intent and actual classroom practice by investigating the human actors navigating these rules. Qualitative methods, such as instructor interviews, can bypass the barrier of closed internal systems to understand why the "silent majority" omits explicit syllabus statements. Furthermore, measuring actual outcomes through surveys similar to the Student Learning Experience Questionnaires used by institutions like HKU (QAC, 2024) will determine whether "Encouraging" policies genuinely foster pedagogical innovation or unintentionally lead to superficial learning.

Note on Generative AI use in this project: In addition to what we documented in the Methods section, Generative AI tools (Gemini, Claude, ChatGPT in their latest versions) were used to support to structure the paper, correct grammatical errors, and support code writing for visualizations and analysis.

References

- Basic Law of the Hong Kong Special Administrative Region of the People's Republic of China, Ch. VI (1990). <https://www.basiclaw.gov.hk/en/basiclaw/chapter6.html>
- Bizikova, L. (2023). Policy coherence and food systems transformation. International Institute for Sustainable Development.
<https://www.iisd.org/articles/policy-analysis/policy-coherence-and-food-systems-transformation>
- Columbia Business School. (n.d.). Getting started with generative artificial intelligence.
<https://business.columbia.edu/samberg/ai/ai-basics>
- Columbia Law School. (n.d.). Columbia Law School interim policy on generative AI.
<https://www.law.columbia.edu/academics/registration-services/resources/columbia-law-school-interim-policy-generative-ai>
- Columbia University Center for Teaching and Learning. (n.d.). Developing AI syllabus statements and communicating with students about AI. <https://ai.ctl.columbia.edu/explore/ai-syllabus-statement/>
- Columbia University Department of Computer Science. (2024). COMS 4705: Natural language processing. <https://www.cs.columbia.edu/~johnhew/coms4705/index.html>
- Columbia University Department of Computer Science. (2025). COMS 3261: Computer science theory.
<https://www.cs.columbia.edu/~josh/cs-theory/>
- Columbia University Department of Psychology. (2024a). PSYC UN 3620: Developmental psychopathology syllabus.
<https://psychology.columbia.edu/sites/psychology.columbia.edu/files/content/PSYC%203620%20Syllabus%20Fall2024.pdf>
- Columbia University Department of Psychology. (2024b). PSYC UN 2235: Thinking and decision making syllabus.
<https://psychology.columbia.edu/sites/psychology.columbia.edu/files/content/UN2235%20Syllabus%20for%20Spring%202024%20KFG.pdf>
- Columbia University Department of Psychology. (2024c). NSBV BC3392: Psychobiology of stress syllabus.
<https://psychology.columbia.edu/sites/psychology.columbia.edu/files/content/Syllabus%20Psychobiology%20of%20Stress%20FA24.pdf>
- Columbia University Information Technology. (n.d.). AI services.
<https://www.cuit.columbia.edu/content/ai-services>
- Columbia University Information Technology. (2025). CUIT AI framework policy.
[https://www.cuit.columbia.edu/sites/www.cuit.columbia.edu/files/content/CUIT%20Website%20Pages/AI%20Framework%20%20\(1\).pdf](https://www.cuit.columbia.edu/sites/www.cuit.columbia.edu/files/content/CUIT%20Website%20Pages/AI%20Framework%20%20(1).pdf)

- Columbia University Information Technology. (2026, January 9). Best practices for responsible AI use.
<https://etc.cuit.columbia.edu/news/best-practices-responsible-ai-use>
- Columbia University Irving Medical Center. (n.d.). AI and generative technology use at CUIMC.
<https://www.it.cuimc.columbia.edu/ai-and-generative-technology-use-cuimc>
- Columbia University Office of the Provost. (2026). Generative AI policy & guidance.
<https://provost.columbia.edu/content/generative-ai>
- Constitutional and Mainland Affairs Bureau. (2018). Guangdong-Hong Kong-Macao Greater Bay Area – Home. The Government of the Hong Kong Special Administrative Region.
<https://www.bayarea.gov.hk/en/home/index.html>
- Digital Policy Office. (2025). Ethical artificial intelligence framework (Version 2.0). The Government of the Hong Kong Special Administrative Region.
https://www.digitalpolicy.gov.hk/en/our_work/data_governance/policies_standards/ethical_ai_framework/doc/Ethical_AI_Framework_en.pdf
- Education Bureau. (2023). Hong Kong's education at a glance. The Government of the Hong Kong Special Administrative Region.
https://www.edb.gov.hk/attachment/en/about-edb/publications-stat/publications-and-videos/HK_Edu_at_a_Glance_EN.pdf
- Jacomy, M., Venturini, T., Heymann, S., & Bastian, M. (2014). ForceAtlas2, a continuous graph layout algorithm for handy network visualization designed for the Gephi software. *PloS one*, 9(6), e98679.
- Loo, B. P. Y. (2025). GEOG3422 [Course outline]. Department of Geography, The University of Hong Kong.
<https://bpyloo.hku.hk/wp-content/uploads/2025/09/GEOG3422-Course-Outline-01Sep2025.pdf>
- Lu, W. (2025, June 20). Hong Kong fast-tracks hotel conversions to narrow student housing gap. *China Daily*. <https://www.chinadailyhk.com/hk/article/614351>
- Ministry of Education. (2026). 教育部等五部门关于印发《"人工智能+教育"行动计划》的通知 ["AI+ Education" action plan]. National Data Administration.
https://www.nda.gov.cn/sjj/zwgk/zcfb/0413/20260410162428109058250_pc.html
- New York City Economic Development Corporation. (2024). NYCEDC releases new report highlighting higher education's critical role in NYC's economy.
<https://edc.nyc/press-release/nycedc-releases-new-report-highlighting-higher-educations-critical-role-nycs-economy>
- New York State Education Department. (2025). College and university evaluation.
<https://www.nysed.gov/college-university-evaluation>

- Newman, M. E., & Girvan, M. (2004). Finding and evaluating community structure in networks. *Physical review E*, 69(2), 026113.
- OECD. (2023). *OECD Digital Education Outlook 2023: Towards an effective digital education ecosystem*. OECD Publishing. <https://doi.org/10.1787/c74f03de-en>
- Peterson, D., Goode, K., & Gehlhaus, D. (2021, September). *Education in China and the United States: A comparative system overview*. Center for Security and Emerging Technology, Georgetown University. <https://cset.georgetown.edu/publication/education-in-china-and-the-united-states/>
- QS Quacquarelli Symonds. (2026). *QS World University Rankings. Top Universities*. <https://www.topuniversities.com/world-university-rankings>
- Quality Assurance Council. (2024). *Report of quality audit of The University of Hong Kong*. University Grants Committee. <https://www.ugc.edu.hk/doc/eng/qac/report/hku202409e.pdf>
- Ruiz, P., Prado, Y., Roschelle, J., & Fusco, J. (2025). *What states say about evaluating AI in education: Reviewing guidance from 32 states and Puerto Rico*. Digital Promise. <https://files.eric.ed.gov/fulltext/ED678842.pdf>
- Sandström, A., Söderberg, C., Lundmark, C., Nilsson, J., & Fjellborg, D. (2020). Assessing and explaining policy coherence: A comparative study of water governance and large carnivore governance in Sweden. *Environmental Policy and Governance*, 30(1), 3–13. <https://doi.org/10.1002/eet.1871>
- Teachers College, Columbia University. (2025). *Recommendations for generative AI use at TC*. <https://www.tc.columbia.edu/digitalfuturestinstitute/ai-in-education/recommendations-for-generative-ai-use-at-tc/>
- Teaching and Learning Innovation Centre, The University of Hong Kong. 2023. "Policy for AI in Education." *AI Education@HKU*. <https://aied.talic.hku.hk/aipolicy/>.
- Teaching and Learning Innovation Centre, The University of Hong Kong. (2026). *Academic integrity and university-wide core principles for using AI in teaching and assessment*. <https://talic.hku.hk/innovation/ai-ethics/core-principles/>
- The Government of the Hong Kong Special Administrative Region. (2025). *The Chief Executive's 2025 Policy Address supplement*. https://www.policyaddress.gov.hk/2025/public/pdf/supplement/supplement-07_en.pdf
- The University of Hong Kong. (2025). *PSYC7319 [Course outline]*. Department of Psychology. <https://psychology.hku.hk/wp-content/programmes/MSocScPsy/CourseOutlines/7319.pdf>
- University Grants Committee. (2025). *UGC Pamphlet 2025. The Government of the Hong Kong Special Administrative Region*. https://www.ugc.edu.hk/doc/eng/ugc/publication/pamphlets/UGC_Pamphlet_2025.pdf

- U.S. Department of Education. (2025, July 22). U.S. Department of Education issues guidance on artificial intelligence use in schools, proposes additional supplemental priority.
<https://www.ed.gov/about/news/press-release/us-department-of-education-issues-guidance-artificial-intelligence-use-schools-proposes-additional-supplemental-priority>
- Wang, H., Dang, A., Wu, Z., & Mac, S. (2024). Generative AI in higher education: Seeing ChatGPT through universities' policies, resources, and guidelines. *Computers and Education: Artificial Intelligence*, 7, 100326. <https://doi.org/10.1016/j.caeai.2024.100326>
- The White House. (2025, April 23). Advancing artificial intelligence education for American youth.
<https://www.whitehouse.gov/presidential-actions/2025/04/advancing-artificial-intelligence-education-for-american-youth/>

Appendix

Appendix A. Coding Rubrics for Dimension

Table A. Policy Stance

Guiding Question: What is the document's overall position on student use of GenAI?

Option	Definition
Prohibitive	Use is viewed as a threat and explicitly banned.
Cautious	Use is a risk and allowed only with permission or strict rules.
Functional	Use is accepted or recognized as a tool (like a calculator), focused on the use of AI and its function rather than whether it is good or bad.
Encouraging	Use is promoted as a benefit to learning or a partner

Table B. Disclosure

Guiding Question: Does the document require students to acknowledge or declare their use of GenAI?

Option	Definition
Required	Students must cite, acknowledge, or declare GenAI use.
Recommended	Disclosure is suggested but not required.
Conditional	Required only for specific tasks (e.g., more than 10% text).
Not specified	No statement about disclosure.

Table C. Discretion

Guiding Question: Who holds the authority to set the rules around GenAI use?

Option	Definition
--------	------------

Gov-mandated	Issued by a government body external to the university
Institution mandated	Issued by a university-wide body (provost, senate, registry)
Faculty/Dept decides	Issued by a sub-unit of the university (school, department)
Instructor decides	Each instructor sets the rule at the course level. Individual course syllabus or instructor statement
Hybrid	The institution establishes a baseline policy (floor) that all must follow, while explicitly delegating additional or more specific rules to faculty or instructors at the course level. Both levels of governance are active simultaneously.
Not specified	Resource guide with no enforceable rule.

Table D. Focus

Guiding Question: What is the primary concern or goal driving the document's GenAI guidance?

Option	Primary Driver	Example
Academic Integrity	Authorship	"Ensuring work is your own original thought..."
AI Literacy	Criticality	"Learning to judge AI biases and hallucinations..."
Pedagogy/Learning	Process	"Using AI to scaffold the brainstorming stage..."
Professional Readiness	Career	"Preparing you for an AI-integrated workplace..."
Privacy/Data/Ethics	Safety	"Protecting your data and avoiding biased tools..."
Equity	Fairness	"Ensuring all students have access to Pro tools..."

Table E. Enforcement

Guiding Question: Does the document specify what happens if GenAI rules are violated?

Primary Driver	Definition
Explicit penalties	A specific consequence for non-compliance is named directly in the document. The reader does not need to consult another document to understand what happens if they violate the rule.
Referenced to integrity code	No penalty is stated directly. The document points to an existing integrity policy or university regulation where consequences are defined.
Guidance only	Recommendations or best practices with no connection to a penalty or integrity code. Non-compliance carries no stated formal consequence.
None stated	The document addresses GenAI but contains no reference to enforcement, penalties, integrity codes, or compliance of any kind. Includes purely descriptive, resource-based, or aspirational documents.

Table F. Coverage

Guiding Question: Who is this document written for and whose behavior does it govern?

Option	Definition
Students only	The document addresses student behavior and obligations exclusively. Instructors appear only as enforcers, not as a governed audience.
Faculty only	The document addresses instructor decisions and practice exclusively. Students appear only as the subjects of those decisions, not as a governed audience.
Students and Faculty	The document explicitly addresses both instructors and students as governed audiences. Both groups have obligations or guidance directed at them.
Researchers only	The document addresses research practice exclusively. Teaching, learning, and student behavior are outside its scope.
Whole	The document addresses all institutional or societal members including

community	administrative staff, researchers, and the general public. Not only for just the teaching and learning relationship.
------------------	--